

An Explicit Decode-Mixing Loss for Mixing-Equivariant Audio Autoencoders

Nurdaulet Akhanov

Mohamed bin Zayed University of Artificial Intelligence

Abu Dhabi, United Arab Emirates

nurdaulet.akhanov@mbzuai.ac.ae

Abstract—Audio mixing is linear in the waveform domain, yet the latent spaces of neural audio autoencoders do not preserve this structure: interpolating two latent codes and decoding does not reproduce the corresponding waveform mix. We introduce an explicit *decode-mixing loss* that enforces this property directly — decoded latent interpolations are required to match true waveform mixes — at the cost of a single extra decoder pass per step. On a waveform GAN autoencoder the loss makes the latent space mixing-equivariant *at no reconstruction cost*, enabling latent-domain stem manipulation: removing a stem by latent subtraction improves from +3.4 to +5.7 dB SI-SDR. At $2\times$ tighter compression the loss is what makes the operation viable at all — against a matched from-scratch control, stem removal goes from an unusable +1.2 dB to +5.2 dB. A controlled ablation isolates the mechanism: the decode-mixing loss itself is the active ingredient, and an adversarial term applied to the mixed path — which inflates a decode-consistency metric — adds no measurable gain on ground-truth equivariance. The recipe transfers across compression rates (trained from scratch at $15.3\times$) and to a structurally different consistency-model autoencoder (Music2Latent), where a continued-training control confirms the gain is the loss and not domain adaptation. Finally, we show that the commonly used decode-vs-decode linearity metric is confounded by decoder phase variance and report a ground-truth-referenced variant throughout.

Index Terms—audio autoencoders, representation learning, mixing equivariance, latent arithmetic, source separation

I. INTRODUCTION

Audio mixing is a linear operation: $\text{mix}(x_1, x_2, \alpha) = \alpha x_1 + (1-\alpha)x_2$. For an autoencoder with encoder f and decoder g , we call the latent space *mixing-equivariant* when

$$g(\alpha f(x_1) + (1-\alpha)f(x_2)) \approx \alpha x_1 + (1-\alpha)x_2. \quad (1)$$

This property turns latent vector arithmetic into audio editing: subtracting a stem’s latent removes it from a mixture, adding latents mixes sources, and weighted combinations remix — all without a dedicated separation or synthesis model. It is also a prerequisite for latent-diffusion systems [1], [2] that compose independently modeled sources. Yet standard autoencoders do not satisfy Eq. 1: their objectives reward reconstruction fidelity, not algebraic structure, and quantized codecs (DAC [3], EnCodec [4]) cannot represent interpolated codes at all.

Recent work pursues linearity through *implicit* regularization — training on random mixes as data augmentation [5] (related to mixup [6]). We take the direct route: an explicit loss (Section III) that compares decoded latent interpolations

to true waveform mixes, propagating gradients through the decoder *and* back into the encoder. Our contributions are:

- 1) An explicit decode-mixing loss that makes a waveform GAN autoencoder mixing-equivariant *without degrading reconstruction* (Section IV-A), enabling latent-arithmetic stem removal (Section IV-C).
- 2) A controlled ablation isolating the mechanism: the decode-mixing loss is the active ingredient, while an adversarial-on-mix term — which inflates a decode-consistency metric — gives no measurable ground-truth gain, and an encoder-only variant and a frozen-encoder variant both underperform.
- 3) Evidence that the recipe generalizes, each against a matched control: from-scratch training at $2\times$ tighter compression — where the loss is what makes latent subtraction viable at all — and transfer to a consistency-model autoencoder with a continued-training control ruling out domain adaptation. The gain holds on all three test domains.
- 4) A methodological correction: the widely used decode-vs-decode SI-SDR_{lin} metric is confounded by decoder phase variance; we report a ground-truth-referenced variant.

II. RELATED WORK

Neural audio codecs. DAC [3], EnCodec [4], and SoundStream [7] pair convolutional encoders with HiFi-GAN decoders [8] and multi-scale discriminators, optimizing reconstruction without latent structure; their residual vector quantization also precludes latent interpolation. RAVE [9] learns a variational latent space but does not target mixing equivariance, and disentanglement regularizers [10] encourage generic smoothness rather than a specific algebraic identity. **Linear audio representations.** Music2Latent [11] is a consistency autoencoder with partial linearity as a byproduct. Torres et al. [5] improve linearity on it via data augmentation. We differ by (i) using an explicit, controllable loss, (ii) isolating which training signal actually drives the property, and (iii) reporting a ground-truth-referenced metric in place of the phase-confounded decode-vs-decode one.

III. METHOD

A. Decode-mixing loss

Given a batch $\{x_i\}$ we form pairs (x_i, x_j) with a fixed-point-free permutation and draw $\alpha \sim U[0, 1]$. Let $\bar{x} = \alpha x_i + (1-\alpha)x_j$ be the waveform mix and $\bar{z} = \alpha f(x_i) + (1-\alpha)f(x_j)$ the latent interpolation. The decode-mixing loss requires the decoded interpolation to match the true mix:

$$\mathcal{L}_{\text{mix}}^{\text{dec}} = \mathcal{L}_{\text{recon}}(g(\bar{z}), \bar{x}), \quad (2)$$

where $\mathcal{L}_{\text{recon}}$ is the multi-resolution STFT + mel loss also used for reconstruction. Gradients flow through g and back into f via $\bar{z}$, so a single term constrains both networks; it costs one extra decoder pass per step. We also study an *encoder-only* variant that penalizes encoder nonlinearity directly, $\mathcal{L}_{\text{mix}}^{\text{enc}} = \|f(\bar{x}) - \bar{z}\|^2$ (one extra encoder pass, no decoder pass). The total generator objective is $\mathcal{L} = \mathcal{L}_{\text{recon}} + \mathcal{L}_{\text{GAN}} + \lambda \mathcal{L}_{\text{mix}}^{\text{dec}} + \lambda \ell_2 \|z\|^2$.

B. Metrics

We report reconstruction $\text{SI-SDR}_{\text{rec}} = \text{SISDR}(g(f(x)), x)$ and two equivariance views at $\alpha=0.5$:

- **SI-SDR_{lin}** (decode-vs-decode): $\text{SISDR}(g(\bar{z}), g(f(\bar{x})))$ — equivariance relative to the model’s *own* reconstruction of the mix.
- **SI-SDR_{lin}^{gt}** (vs. ground truth): $\text{SISDR}(g(\bar{z}), \bar{x})$ — the quantity in Eq. 1, comparable across models.

We also report $\ell_{\text{lat}} = \|\bar{z} - f(\bar{x})\|^2 / \|f(\bar{x})\|^2$ (encoder linearity) and $\text{MixRate} = \mathcal{L}_{\text{recon}}(g(\bar{z}), \bar{x}) / \mathcal{L}_{\text{recon}}(g(f(\bar{x})), \bar{x})$ (decoder equivariance; < 1 is better than the encode-then-decode path). $\text{SI-SDR}_{\text{lin}}$ is the metric used by prior work, but it is *gameable*: a decoder driven onto a common realistic manifold raises the mutual SI-SDR of its two decode paths without bringing either closer to $\bar{x}$ (Section IV-A). It is also confounded for stochastic decoders: a consistency-model decoder [11] draws fresh noise per call, so two decodes of the same latent differ in phase and $\text{SI-SDR}_{\text{lin}}$ collapses unless the noise is shared. Fig. 1 shows the same Music2Latent checkpoints scoring ≈ -8 dB under independent-noise decoding (the protocol implied by prior reports) and *positive* under shared-noise decoding — a 12 dB swing driven purely by the eval protocol, not the model. We therefore treat $\text{SI-SDR}_{\text{lin}}^{\text{gt}}$ and MixRate as the equivariance metrics of record, and use shared decode noise everywhere.

C. Experimental setup

The primary model is a waveform autoencoder (DAC-style 1-D convolutional encoder, HiFi-GAN decoder [8], combined multi-scale-STFT and multi-period discriminator) operating at 44.1 kHz on 1-second mono chunks, trained on a ~ 950 -hour union of FMA-large, MAESTRO, and MUSDB18-HQ [12].¹ Two compression regimes are used: $7.66\times$ ($d=128$) and $15.3\times$ ($d=64$). Mixing losses are added either as a 25k-step fine-tune over a converged baseline (v2.x) or from scratch over 250k steps (v3.x). All runs use a single NVIDIA RTX 5090.

¹Code, training configurations, pretrained weights, and every metric reported here: <https://github.com/nurdauletakhanov/musicgen>




Fig. 1. Decode-noise protocol confound (Music2Latent, MUSDB18, $\alpha=0.5$). The same checkpoints score deeply negative under independent-noise decoding and positive under shared-noise decoding; the ≈ -8 dB figure is a phase artifact, not a linearity measurement.




Fig. 2. Decode-mixing loss: encode a pair, interpolate latents, decode, and compare to the true waveform mix. Gradients (dashed) flow through the decoder and back into the encoder via $\bar{z}$.

Unless noted, metrics are computed over the full multi-source test set; the α -sweep (Section IV-D) uses a source-balanced subsample.

IV. RESULTS

A. Mechanism ablation

Table I ablates the loss on the $7.66\times$ GAN autoencoder (25k-step fine-tunes from a shared converged baseline). Three observations drive the rest of the paper.

(1) **Reconstruction is unaffected.** SDR_{rec} is flat across all configurations (11.2–11.4 dB), and so is FAD (0.044–0.046): the mixing loss buys equivariance for free, on both a sample-wise and a distributional measure of quality.

(2) **The decode-mixing loss is the active ingredient; the discriminator on the mixed path is not.** On the ground-truth metric, $\mathcal{L}_{\text{mix}}^{\text{dec}}$ alone improves $\text{SDR}_{\text{lin}}^{\text{gt}}$ from 8.0 to 10.2 dB and is the only run with $\text{MixRate} < 1$ (0.97). Adding an adversarial term to the mixed path (“+disc-on-mix”) does *not* improve ground-truth equivariance (9.9 dB, MixRate 1.05) — yet it dramatically inflates the decode-vs-decode SDR_{lin} (12.1 $\rightarrow$ 16.0 dB). The discriminator pulls the two decode paths onto a common realistic manifold, raising their mutual SI-SDR without moving either toward $\bar{x}$. This is direct evidence that SDR_{lin} is gameable and that no adversarial machinery is required for equivariance.

(3) **Direct decoder gradient matters.** The encoder-only variant $\mathcal{L}_{\text{mix}}^{\text{enc}}$ drives ℓ_{lat} lowest but leaves MixRate high (1.15–1.18): it linearizes the encoder without making the decoder

Loss	SDR _{rec}	SDR _{lin}	SDR _{lin} ^{gt}	ℓ_{lat}	MixRate	FAD
baseline (no mix)	+11.3	+12.0	+8.0	0.068	1.203	0.045
$\mathcal{L}_{\text{dec}}$	+11.4	+12.1	+10.2	0.052	0.972	0.044
$\mathcal{L}_{\text{dec}}$ + disc-on-mix	+11.4	+16.0	+9.9	0.050	1.052	0.044
$\mathcal{L}_{\text{enc}}$ ($\gamma=5$)	+11.3	+12.8	+8.4	0.044	1.175	0.044
$\mathcal{L}_{\text{enc}}$ ($\gamma=10$)	+11.3	+13.1	+8.5	0.039	1.165	0.044
$\mathcal{L}_{\text{enc}}$ ($\gamma=20$)	+11.3	+13.5	+8.7	0.034	1.153	0.045
$\mathcal{L}_{\text{dec}}$ frozen-enc	+11.2	+11.7	+9.5	0.072	1.056	0.046

TABLE I

MECHANISM ABLATION ON THE GAN AUTOENCODER (7.66 $\times$, FULL TEST SET, $\alpha=0.5$). SDR_{lin} IS DECODE-VS-DECODE; SDR_{lin}^{gt} IS VS. GROUND TRUTH. SINGLE SEED.

Model	SDR _{rec}	SDR _{lin}	SDR _{lin} ^{gt}	MixRate	FAD
GAN AE, 7.66 $\times$, no mix	+11.3	+12.0	+8.0	1.203	0.045
GAN AE, 7.66 $\times$, +mix	+11.4	+12.1	+10.2	0.972	0.044
GAN AE, 15.3 $\times$, no mix	+10.0	+10.6	+6.8	1.187	0.052
GAN AE, 15.3 $\times$, +mix	+9.9	+14.6	+8.6	1.046	0.061
M2L, no mix (baseline)	-3.3	+5.0	-9.6	1.101	0.127
M2L, fine-tune control	-3.1	+5.1	-9.5	1.125	—
M2L, +mix	-2.6	+6.9	-7.8	1.086	0.119

TABLE II

GENERALIZATION ACROSS COMPRESSION AND ARCHITECTURE (FULL TEST SET, $\alpha=0.5$, EMA WEIGHTS FOR MUSIC2LATENT). THE M2L “FINE-TUNE CONTROL” IS THE CONTINUED-TRAINING BASELINE WITH MIXING DISABLED.

equivariant. Freezing the encoder and training only the decoder is the single worst run on the ground-truth metric relative to its reconstruction, confirming both networks must adapt.

B. Generalization: compression and architecture

Table II extends the recipe along two axes. **Compression:** trained *from scratch* at 15.3 $\times$ (v3.1), the mixing model matches its no-mixing control (v3.0) on reconstruction (9.9 vs. 10.0 dB) while improving ground-truth equivariance by +1.7 dB (SDR_{lin}^{gt} 6.8 $\rightarrow$ 8.6), MixRate (1.19 $\rightarrow$ 1.05), and encoder linearity (ℓ_{lat} 0.078 $\rightarrow$ 0.057). Because v3.0 is a matched control — same architecture, compression, and schedule, mixing disabled — this isolates the loss’s contribution and shows the effect is not an artifact of the looser 7.66 $\times$ rate or of warm-starting. **Architecture:** we fine-tune the published Music2Latent [11] consistency autoencoder with a decode-mixing objective (its native consistency loss applied to the mixed pair substitutes for the discriminator). A continued-training control — the same fine-tune with the mixing terms disabled — moves equivariance by < 0.2 dB and *worsens* FAD, so the mixing model’s gain is attributable to the loss, not to domain adaptation. Both decode-vs-decode and ground-truth metrics agree on the improvement, confirming it is not a measurement artifact. Absolute SDR_{lin}^{gt} remains negative on Music2Latent because its consistency decoder is phase-incoherent with raw waveforms; the latent *geometry* improves (MixRate, ℓ_{lat}) but a consistency decoder does not convert this into clean waveform arithmetic as the GAN decoder does — hence we report the GAN autoencoder as the primary result.

The gain is not domain-driven. Our test set spans three acoustically distinct domains — crowd-sourced indie produc-

SDR _{lin} ^{gt}	FMA	MAESTRO	MUSDB	all
7.66 $\times$, no mix	+6.8	+13.8	+6.2	+8.0
7.66 $\times$, +mix	+8.9	+16.9	+8.2	+10.2
Δ	+2.1	+3.1	+2.0	+2.3
15.3 $\times$, no mix	+5.8	+12.0	+5.5	+6.8
15.3 $\times$, +mix	+7.2	+15.1	+6.7	+8.6
Δ	+1.4	+3.2	+1.2	+1.7

TABLE III

GROUND-TRUTH EQUIVARIANCE PER TEST DOMAIN (dB). EACH Δ IS AGAINST THE MATCHED NO-MIXING CONTROL AT THE SAME COMPRESSION RATE. THE GAIN IS PRESENT ON EVERY DOMAIN, NOT CARRIED BY ONE.

tion (FMA), solo piano (MAESTRO), and multi-track studio mixes (MUSDB18). Table III breaks SDR_{lin}^{gt} out by domain against the matched no-mixing control at each compression rate. The improvement appears on all three at both rates (+2.0 to +3.1 dB at 7.66 $\times$; +1.2 to +3.2 dB at 15.3 $\times$), so it reflects a property of the learned latent space rather than an artifact of one recording style. The absolute ordering (MAESTRO easiest, MUSDB hardest) tracks reconstruction difficulty and is preserved by the mixing loss.

C. Downstream: latent-arithmetic stem removal

We remove a stem by latent subtraction, $\hat{x} = g(f(x_{\text{mix}}) - f(x_{\text{stem}}))$, and measure SI-SDR against the true residual on the MUSDB18 test set. This tests *extrapolation* beyond the convex range $\alpha \in [0, 1]$ seen in training. Alongside the subtraction score we report the *linearity tax* gap = SISDR($g(f(x_{\text{res}})), x_{\text{res}}$) - SISDR($\hat{x}, x_{\text{res}}$), the shortfall of latent arithmetic against the same model’s own encode-decode ceiling on the residual. The gap normalizes away raw reconstruction quality and so compares models at different compression rates on equal terms.

Table IV shows the mixing loss roughly doubles usable subtraction quality at 7.66 $\times$ (+3.4 $\rightarrow$ +5.7 dB overall) and cuts the linearity tax from 5.2 to 3.0 dB, with the largest gain on bass (+0.4 $\rightarrow$ +4.9 dB). Adding disc-on-mix is comparable rather than better (+6.0 dB, tax 2.6), consistent with Section IV-A. The effect is starkest against the matched from-scratch control at 15.3 $\times$: without the mixing loss, latent subtraction is essentially unusable (+1.2 dB overall, and -1.6 dB on bass — worse than emitting the unmodified mixture), while the same architecture trained with it reaches +5.2 dB and the *lowest* tax of any model (2.5 dB). Tighter compression therefore does not merely cost a little subtraction quality — it is what makes the property fail outright, and the loss is what restores it.

D. Equivariance across mixing ratios

Sweeping α on a source-balanced subsample (Fig. 3), the mixing loss *flattens* the equivariance curve: the no-mixing baseline sags most at the hardest balanced mix ($\alpha=0.5$) and recovers toward the extremes, while mixing models stay near-uniform across α (e.g. SDR_{lin}^{gt} swing of 1.2 dB for $\mathcal{L}_{\text{mix}}^{\text{dec}}$ + disc vs. 3.2 dB for the baseline). Thus the property holds across mixing ratios, not only at the trained operating point.

Model	drums	bass	vocals	other	all	gap↓
7.66× no mix	+4.2	+0.4	+5.2	+3.9	+3.4	5.2
7.66× + $\mathcal{L}_{\text{dec}}$	+5.7	+4.9	+7.2	+5.1	+5.7	3.0
7.66× + $\mathcal{L}_{\text{dec}}$ +disc	+6.2	+5.4	+7.5	+5.0	+6.0	2.6
15.3× no mix	+1.2	-1.6	+3.0	+2.1	+1.2	6.6
15.3× +mix	+5.4	+4.7	+6.9	+3.9	+5.2	2.5
M2L +mix	-9.9	-12.3	-9.9	-12.7	-11.2	8.4

TABLE IV

STEM REMOVAL VIA LATENT SUBTRACTION (SI-SDR, dB; MUSDB18 TEST). HIGHER IS BETTER; “GAP” IS THE LINEARITY TAX AGAINST EACH MODEL’S OWN ENCODE-DECODE CEILING (LOWER IS BETTER).




Fig. 3. Ground-truth equivariance vs. mixing coefficient α (source-balanced subsample). The no-mixing baseline sags at the hardest balanced mix ($\alpha=0.5$); the decode-mixing loss flattens and lifts the curve across all α .

We caution that the decode-vs-decode SDR_{lin} is U-shaped in α for *all* models — near-identity mixes are trivially easy — which is a property of that metric, distinct from the phase-variance confound discussed in Section III-B.

V. CONCLUSION

An explicit decode-mixing loss makes a waveform autoencoder mixing-equivariant at no reconstruction cost, enabling latent-arithmetic audio editing. The loss itself — not an adversarial term on the mixed path — is the active ingredient, and the effect transfers across compression rates and to a consistency-model autoencoder under a continued-training control. **Limitations.** All results use a single training seed; small differences (e.g. $\mathcal{L}_{\text{mix}}^{\text{dec}}$ vs. $\mathcal{L}_{\text{mix}}^{\text{dec}} + \text{disc}$) are within plausible seed variance, and we report them as “comparable.” Experiments are mono; consistency-decoder latent arithmetic remains phase-limited; and we conduct no listening test. Stereo, multi-seed estimates, and latent-diffusion composition are future work.

REFERENCES

- [1] H. Liu, Z. Chen, Y. Yuan, X. Mei, X. Liu, D. Mandic, W. Wang, and M. D. Plumbley, “AudioLDM: Text-to-audio generation with latent diffusion models,” in *Proc. of the 40th Int. Conf. on Machine Learning (ICML)*, 2023.
- [2] Z. Evans, J. Parker, C. Carr, Z. Zukowski, J. Taylor, and J. Pons, “Stable audio open,” in *Proc. of the 25th Int. Society for Music Information Retrieval Conf. (ISMIR)*, 2024.

- [3] R. Kumar, P. Seetharaman, A. Luebs, I. Kumar, and K. Kumar, “High-fidelity audio compression with improved RVQGAN,” in *Proc. of the 38th Int. Conf. on Neural Information Processing Systems (NeurIPS)*, 2024.
- [4] A. Défossez, J. Copet, G. Synnaeve, and Y. Adi, “High fidelity neural audio compression,” *Trans. on Machine Learning Research*, 2023.
- [5] B. Torres, M. Moussallam, and G. Meseguer-Brocal, “Learning linearity in audio consistency autoencoders via implicit regularization,” in *Proc. of the IEEE Int. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, 2026.
- [6] H. Zhang, M. Cisse, Y. N. Dauphin, and D. Lopez-Paz, “mixup: Beyond empirical risk minimization,” in *Proc. of the 6th Int. Conf. on Learning Representations (ICLR)*, 2018.
- [7] N. Zeghidour, A. Luebs, A. Omran, J. Skoglund, and M. Tagliasacchi, “SoundStream: An end-to-end neural audio codec,” *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, vol. 30, pp. 2709–2723, 2022.
- [8] J. Kong, J. Kim, and J. Bae, “HiFi-GAN: Generative adversarial networks for efficient and high fidelity speech synthesis,” in *Proc. of the 34th Int. Conf. on Neural Information Processing Systems (NeurIPS)*, 2020.
- [9] A. Caillon and P. Esling, “RAVE: A variational autoencoder for fast and high-quality neural audio synthesis,” *arXiv preprint arXiv:2111.05011*, 2021.
- [10] I. Higgins, L. Matthey, A. Pal, C. Burgess, X. Glorot, M. Botvinick, S. Mohamed, and A. Lerchner, “ β -VAE: Learning basic visual concepts with a constrained variational framework,” in *Proc. of the 5th Int. Conf. on Learning Representations (ICLR)*, 2017.
- [11] M. Pasini and J. Schlüter, “Music2Latent: Consistency autoencoders for latent audio compression,” in *Proc. of the 41st Int. Conf. on Machine Learning (ICML)*, 2024.
- [12] Z. Rafii, A. Liutkus, F.-R. Stöter, S. I. Mimilakis, and R. Bittner, “MUSDB18-HQ — an uncompressed version of MUSDB18,” <https://zenodo.org/record/3338373>, 2019.