

What Makes Audio Latents Mixing-Equivariant?

A Controlled Study of Explicit Supervision

Nurdaulet Akhanov

Mohamed bin Zayed University of Artificial Intelligence

Abu Dhabi, United Arab Emirates

nurdaulet.akhanov@mbzuai.ac.ae

Abstract—Mixing is linear on waveforms, so a latent space in which interpolating two codes and decoding reproduces the waveform mix is an *equivariant* representation. Recent work induces this by decoding averaged latents against the true mix. We ask which part of that supervision is responsible and whether the certifying metrics measure it at all. An explicit decode-mixing loss suffices: on a waveform GAN autoencoder it makes the latent space mixing-equivariant with reconstruction preserved, and it transfers to $2\times$ tighter compression. An adversarial term on the mixed path adds no ground-truth gain while inflating a decode-consistency metric. Latent stem subtraction, in contrast, is governed largely by the latent origin: adding the encoded silence $f(0)$ reduces the loss’s apparent subtraction advantage by 87–95% on a pre-registered 240-track held-out corpus. Comparisons are single-run and matched at one initialization.

Index Terms—representation learning, equivariant representations, audio autoencoders, latent arithmetic, controlled ablation

I. INTRODUCTION

Audio mixing is a linear operation: $\text{mix}(x_1, x_2, \alpha) = \alpha x_1 + (1-\alpha)x_2$. For an autoencoder with encoder f and decoder g , we call the latent space *mixing-equivariant* when

$$g(\alpha f(x_1) + (1-\alpha)f(x_2)) \approx \alpha x_1 + (1-\alpha)x_2. \quad (1)$$

Eq. 1 constrains combinations whose coefficients sum to one, so latent *subtraction* $g(f(x_{\text{mix}}) - f(x_{\text{stem}}))$, with coefficients $(1, -1)$, is not implied by it: an offset $f(x) = x + b$, $g(z) = z - b$ satisfies Eq. 1 exactly yet subtracts to $x_{\text{res}} - b$. For an affine $f(x) = Ax + b$, adding the encoded silence restores the sum to one and $f(x_{\text{mix}}) - f(x_{\text{stem}}) + f(0) = f(x_{\text{res}})$, so the decode is the model’s own reconstruction of the residual, whatever the decoder; matching the residual waveform additionally requires exact reconstruction. How far real encoders depart from affine is empirical (Section IV-C). Where it holds, latent arithmetic becomes audio editing without a dedicated separation or synthesis model, and it is a prerequisite for latent-diffusion systems [1], [2] that compose independently modeled sources. Yet standard autoencoders do not satisfy Eq. 1: their objectives reward reconstruction fidelity, not algebraic structure, and in quantized codecs (DAC [3], EnCodec [4]) an interpolated point is not a codebook entry, so code-domain mixing is undefined without leaving the quantizer.

Torres et al. [5] induce linearity in a consistency autoencoder by decoding the *average* of separately encoded latents and training it to reconstruct the true mix, a supervision

principle related to mixup [6]. This leaves open which part of that supervision is responsible for the property, whether the encoder or the decoder must carry it, whether a consistency decoder can turn latent linearity into usable waveform arithmetic, and whether the metrics used to certify it measure it at all. We answer these questions with a controlled study: the same principle written as an explicit loss (Section III) on a waveform GAN autoencoder, whose effect we then isolate by intervening on one factor at a time. Our contributions are:

- 1) A controlled study of mixed-latent supervision on a waveform GAN autoencoder: the decode-mixing loss alone is sufficient, an adversarial term on the mixed path adds no ground-truth gain on the tested contrast, and decoder-only adaptation retains most of the improvement while joint adaptation is best among the tested configurations (Section IV-A).
- 2) Generalization against matched controls: the combined $\mathcal{L}_{\text{mix}}^{\text{dec}} + \text{disc}$ recipe trained from scratch at $2\times$ tighter compression, and a re-run on a consistency-model autoencoder where latent geometry improves but waveform arithmetic does not (Section IV-B).
- 3) A downstream analysis showing latent stem removal is limited by the latent origin rather than the loss, for every evaluated GAN configuration (Section IV-C).
- 4) A methodological correction: the widely used decode-vs-decode SI- SDR_{lin} is sensitive to the decode-noise protocol and to an adversarial term; we report a ground-truth-referenced variant.

II. RELATED WORK

Neural audio codecs. DAC [3], EnCodec [4], and SoundStream [7] pair convolutional encoders with HiFi-GAN decoders [8] and multi-scale discriminators, optimizing reconstruction without latent structure; their residual vector quantization also precludes latent interpolation. RAVE [9] learns a variational latent space but does not target mixing equivariance, and disentanglement regularizers [10] encourage generic smoothness rather than a specific algebraic identity. **Equivariant representations.** Mixing-equivariance follows the symmetry-based view of representation learning [11], though convex mixing is an affine combination, not a group action, so we borrow the framing and not the formalism. We treat equivariance as a property to be *induced* by supervision

and isolate which supervision induces it. **Linear audio representations.** Music2Latent [12] is a consistency autoencoder with partial linearity as a byproduct. Torres et al. [5] make it approximately linear by conditioning the decoder on the average of separately encoded latents while denoising the true mix under the unchanged consistency loss (plus a random gain for homogeneity, which we do not study). Our $\mathcal{L}_{\text{mix}}^{\text{dec}}$ is the explicit-loss form of that additivity supervision: the same input (averaged latents) and target (true mix), differing in loss family (multi-resolution STFT and mel versus consistency), mixing coefficient ($\alpha \sim U[0,1]$ versus 0.5), and decoder class (deterministic GAN versus consistency model). We add the analysis: which ingredient carries the effect, a matched re-run on their architecture with a continued-training control (Section IV-B), and ground-truth-referenced metrics. Their additivity (SNR, spectral distance) and separation scores compare two decodes, the protocol Section III-B shows to be confounded, and their separation is negative on every stem (-0.3 to -1.8 dB), consistent with this sampling sensitivity.

III. METHOD

A. Decode-mixing loss

Given a batch $\{x_i\}$ we form pairs (x_i, x_j) with a fixed-point-free permutation and draw $\alpha \sim U[0,1]$. Let $\bar{x} = \alpha x_i + (1-\alpha)x_j$ be the waveform mix and $\bar{z} = \alpha f(x_i) + (1-\alpha)f(x_j)$ the latent interpolation. The decode-mixing loss requires the decoded interpolation to match the true mix:

$$\mathcal{L}_{\text{mix}}^{\text{dec}} = \mathcal{L}_{\text{recon}}(g(\bar{z}), \bar{x}) + \mathcal{L}_{\text{wave-}\ell_1}(g(\bar{z}), \bar{x}), \quad (2)$$

where $\mathcal{L}_{\text{recon}}$ is the weighted multi-resolution STFT and mel loss also used for reconstruction and $\mathcal{L}_{\text{wave-}\ell_1}$ is the unweighted mean absolute error over all samples and batch elements. The reconstruction path carries no waveform term, so the two paths differ in objective. Gradients flow through g and back into f via $\bar{z}$, so a single term constrains both networks; it costs one extra decoder pass per step. We also study an *encoder-only* variant that penalizes encoder non-linearity directly, $\mathcal{L}_{\text{mix}}^{\text{enc}} = \|f(\bar{x}) - \bar{z}\|^2$ (one extra encoder pass, no decoder pass). The total generator objective is $\mathcal{L} = \mathcal{L}_{\text{recon}} + \mathcal{L}_{\text{GAN}} + \lambda \mathcal{L}_{\text{mix}}^{\text{dec}} + \lambda_{\ell_2} \|z\|^2$.

B. Metrics

We report reconstruction $\text{SI-SDR}_{\text{rec}} = \text{SISDR}(g(f(x)), x)$ and two equivariance views at $\alpha=0.5$:

- **SI-SDR_{lin}** (decode-vs-decode): $\text{SISDR}(g(\bar{z}), g(f(\bar{x})))$, equivariance relative to the model’s *own* reconstruction of the mix.
- **SI-SDR_{lin}^{gt}** (vs. ground truth): $\text{SISDR}(g(\bar{z}), \bar{x})$, the quantity in Eq. 1 up to scale, comparable across models. SI-SDR fits a gain before scoring, so it does not assess level preservation; we report the fitted gain separately (Section IV-B). FAD uses music-finetuned LAION-CLAP embeddings of non-overlapping 10 s windows at 48 kHz, not VGGish, so the values are not comparable with VGGish-based FAD; the fitted gain is measured on 3,000 samples.




Fig. 1. Decode-noise protocol sensitivity (Music2Latent, MUSDB18, $\alpha=0.5$). Changing only the coupling of decoder noise changes the score by approximately 13 dB on identical pairs.

We also report $\ell_{\text{lat}} = \|\bar{z} - f(\bar{x})\|^2 / \|f(\bar{x})\|^2$ (encoder linearity) and $\text{MixRate} = \mathcal{L}_{\text{mix}}(g(\bar{z}), \bar{x}) / \mathcal{L}_{\text{mix}}(g(f(\bar{x})), \bar{x})$, with $\mathcal{L}_{\text{mix}}$ the three-term mixed-path loss above in both numerator and denominator, (decoder equivariance; < 1 is better than the encode-then-decode path). SI-SDR_{lin} shares the decode-vs-decode structure of prior work’s decoder-additivity metrics (SNR and spectral distance between two decodes [5]), and the two views can be moved apart: an adversarial term on the mixed path raises the mutual SI-SDR of the two decode paths without bringing either closer to $\bar{x}$ (Section IV-A). We do not identify the mechanism; one candidate is that both decodes are driven onto a common realistic manifold. It is also confounded for stochastic decoders: a consistency-model decoder [12] draws fresh noise per call. Decoding the *same* latent twice is deterministic under shared noise but scores -1.8 dB under independent noise: decoder sampling substantially affects decode-vs-decode scores even when no latent arithmetic is performed. On identical pairs, Fig. 1 shows the same Music2Latent checkpoints scoring -9.1 and -7.3 dB when the two decodes draw independent noise and $+3.6$ and $+5.9$ dB when they share it, a 13 dB swing driven by the evaluation protocol rather than the model. Which protocol prior work used is not stated in enough detail for us to determine. This shows protocol sensitivity; decoding the same latent twice with independent noise isolates how much of the disagreement arises without any latent arithmetic (Section IV-B). We therefore treat SI-SDR_{lin}^{gt} and MixRate as the equivariance metrics of record, and use shared decode noise everywhere.

C. Experimental setup

The primary model is a waveform autoencoder (DAC-style 1-D convolutional encoder, HiFi-GAN decoder [8], combined multi-scale-STFT and multi-period discriminator) operating at 44.1 kHz on 1-second mono chunks, trained on a ~ 950 -hour union of FMA-large, MAESTRO, and MUSDB18-HQ [13].¹

¹Code, training configurations, pretrained weights, every metric reported here, and audio examples: <https://github.com/nurdauletakhanov/musicgen>

Loss	SDR _{rec}	SDR _{lin}	SDR _{lin} ^{gt}	ℓ_{lat}	MixRate	FAD
baseline (no mix)	+11.3	+12.0	+8.0	0.068	1.203	0.045
$\mathcal{L}_{\text{dec}}$	+11.4	+12.1	+10.2	0.052	0.972	0.044
$\mathcal{L}_{\text{dec}} + \text{disc-on-mix}$	+11.4	+16.0	+9.9	0.050	1.052	0.044
$\mathcal{L}_{\text{enc}} (\gamma=5)$	+11.3	+12.8	+8.4	0.044	1.175	0.044
$\mathcal{L}_{\text{enc}} (\gamma=10)$	+11.3	+13.1	+8.5	0.039	1.165	0.044
$\mathcal{L}_{\text{enc}} (\gamma=20)$	+11.3	+13.5	+8.7	0.034	1.153	0.045
$\mathcal{L}_{\text{dec}}$ frozen-enc	+11.2	+11.7	+9.5	0.072	1.056	0.046

TABLE I

MECHANISM ABLATION ON THE GAN AUTOENCODER (7.66 $\times$, FULL TEST SET, $\alpha=0.5$). SDR_{lin} IS DECODE-VS-DECODE; SDR_{lin}^{gt} IS VS. GROUND TRUTH. SINGLE SEED.

Two compression regimes are used: 7.66 $\times$ ($d=128$) and 15.3 $\times$ ($d=64$). Mixing losses are added either as a 25k-step fine-tune over a converged baseline (v2.x) or from scratch over 250k steps (v3.x). **Checkpoint selection.** The released weights are the checkpoint with the lowest validation training loss, and the validation loader read the test split; for every v2.x run this coincides with the final step, so selection did not alter those weights, while both v3.x runs were selected at step 240k of 250k. Contrasts are therefore between matched configurations under identical selection, but absolute generalization claims should be read with this in mind; the code now supports a held-out validation split. All runs use a single NVIDIA RTX 5090. Unless noted, metrics are computed over the full multi-source test set; the α -sweep (Section IV-D) uses a source-balanced subsample.

IV. RESULTS

A. Which signal causes equivariance? An interventional ablation

Table I ablates the loss on the 7.66 $\times$ GAN autoencoder (25k-step fine-tunes from a shared converged baseline). Each row differs from its comparator in exactly one factor: $\mathcal{L}_{\text{mix}}^{\text{dec}}$ and $\mathcal{L}_{\text{mix}}^{\text{enc}}$ against the baseline (the loss), $\mathcal{L}_{\text{mix}}^{\text{dec}} + \text{disc}$ against $\mathcal{L}_{\text{mix}}^{\text{dec}}$ (the adversarial term), and frozen-encoder $\mathcal{L}_{\text{mix}}^{\text{dec}}$ against $\mathcal{L}_{\text{mix}}^{\text{dec}}$ (the encoder’s freedom to adapt); the design is not a full factorial. Three observations drive the paper.

(1) Reconstruction is similar across configurations. SDR_{rec} spans 11.2–11.4 dB and FAD 0.044–0.046, so we observe no reconstruction cost at 7.66 $\times$. These are single runs, so an absence this small is not established beyond seed variance.

(2) The decode-mixing loss is the active ingredient. On the ground-truth metric, $\mathcal{L}_{\text{mix}}^{\text{dec}}$ alone improves SDR_{lin}^{gt} from 8.0 to 10.2 dB and is the only run with MixRate < 1 (0.97). Adding an adversarial term to the mixed path (“+disc-on-mix”) does *not* improve ground-truth equivariance (9.9 dB, MixRate 1.05), yet it dramatically inflates the decode-vs-decode SDR_{lin} (12.1 $\rightarrow$ 16.0 dB). The two decode paths move toward each other without either moving toward $\bar{x}$; one candidate mechanism is a common realistic manifold. This is direct evidence that SDR_{lin} can rise without ground-truth equivariance rising,

Model	SDR _{rec}	SDR _{lin}	SDR _{lin} ^{gt}	MixRate	FAD
GAN, 7.66 $\times$, no mix	+11.3	+12.0	+8.0	1.203	0.045
GAN, 7.66 $\times$, + $\mathcal{L}_{\text{dec}}$	+11.4	+12.1	+10.2	0.972	0.044
GAN, 15.3 $\times$, no mix	+10.0	+10.6	+6.8	1.187	0.052
GAN, 15.3 $\times$, + $\mathcal{L}_{\text{dec}} + \text{disc}$	+9.9	+14.6	+8.6	1.046	0.061
M2L, no mix (baseline)	-3.3	+5.0	-9.6	1.101	0.127
M2L, fine-tune control	-3.1	+5.1	-9.5	1.125	0.134
M2L, +mix	-2.6	+6.9	-7.8	1.086	0.119

TABLE II

GENERALIZATION ACROSS COMPRESSION AND ARCHITECTURE (FULL TEST SET, $\alpha=0.5$, EMA WEIGHTS FOR MUSIC2LATENT). THE M2L “FINE-TUNE CONTROL” IS THE CONTINUED-TRAINING BASELINE WITH MIXING DISABLED.

and that given $\mathcal{L}_{\text{mix}}^{\text{dec}}$ no adversarial term improved the ground-truth metric here; whether the discriminator has an effect on its own, without $\mathcal{L}_{\text{mix}}^{\text{dec}}$, is not tested here.

(3) Decoder-only adaptation retains most of the improvement. The encoder-only variant $\mathcal{L}_{\text{mix}}^{\text{enc}}$ drives ℓ_{lat} lowest but leaves MixRate high (1.15–1.18): it linearizes the encoder without making the decoder equivariant. Training only the decoder under $\mathcal{L}_{\text{mix}}^{\text{dec}}$ retains most of the gain (9.5 dB versus 10.2 joint and 8.0 baseline) while leaving encoder linearity unimproved. Joint training is best among the tested configurations. That decoder adaptation is largely sufficient here does not show it is necessary, which would need a decoder-frozen counterpart we did not train.

B. Intervening on compression and architecture

Table II extends the recipe along two axes. **Compression:** trained *from scratch* at 15.3 $\times$ (v3.1, which uses the combined $\mathcal{L}_{\text{mix}}^{\text{dec}} + \text{disc}$ recipe, so it tests that combination rather than $\mathcal{L}_{\text{mix}}^{\text{dec}}$ alone), the mixing model matches its no-mixing control (v3.0) on reconstruction (9.9 vs. 10.0 dB) while improving ground-truth equivariance by +1.7 dB (SDR_{lin}^{gt} 6.8 $\rightarrow$ 8.6), MixRate (1.19 $\rightarrow$ 1.05), and encoder linearity (ℓ_{lat} 0.078 $\rightarrow$ 0.057). FAD, however, rises from 0.052 to 0.061 (Table II); whether that is perceptible is untested. At 7.66 $\times$ the reported reconstruction metrics are similar in these runs (Table I); at 15.3 $\times$ that similarity holds for SDR_{rec} only. Because v3.0 is a matched control, the effect is not an artifact of the looser 7.66 $\times$ rate or of warm-starting. **Level preservation.** The gain SI-SDR discards: the decoded mix is 1.6 dB too quiet without the loss and 0.9 dB with it, against a reconstruction level error of 0.8 dB. The loss reduces level error as well as spectral error. **Architecture:** fine-tuning the published Music2Latent [12] consistency autoencoder with the same mixed-latent supervision, using its native consistency loss, improves SDR_{lin}^{gt} against a continued-training control (−9.5 $\rightarrow$ −7.8 dB) while that control moves equivariance by < 0.2 dB and *worsens* FAD, so the gain is the loss, not domain adaptation. Latent geometry and ground-truth mixing both improve, yet waveform arithmetic does not become practical. The systems differ in too many respects for this to isolate a cause; we treat it as a limit on transfer.

Model	drums	bass	vocals	other	all	gap↓
7.66× no mix	+4.2	+0.4	+5.2	+4.0	+3.4	5.2
+ $f(0)$	+8.8	+6.8	+9.6	+8.0	+8.3	0.4
7.66× + $\mathcal{L}_{\text{dec}}$	+5.7	+4.9	+7.3	+5.2	+5.8	2.9
+ $f(0)$	+8.5	+6.9	+9.4	+8.3	+8.3	0.5
7.66× + $\mathcal{L}_{\text{dec}}$ +disc	+6.2	+5.4	+7.5	+5.2	+6.1	2.6
+ $f(0)$	+8.5	+7.0	+9.3	+8.1	+8.2	0.4
15.3× no mix	+1.2	-1.6	+3.0	+2.2	+1.2	6.6
+ $f(0)$	+7.3	+5.8	+8.8	+7.2	+7.3	0.5
15.3× + $\mathcal{L}_{\text{dec}}$ +disc	+5.4	+4.7	+6.9	+4.0	+5.2	2.4
+ $f(0)$	+7.3	+6.1	+8.5	+7.2	+7.3	0.4
M2L +mix	-9.9	-12.3	-9.9	-12.7	-11.2	8.4

TABLE III

STEM REMOVAL VIA LATENT SUBTRACTION (SI-SDR, dB; MUSDB18 TEST). HIGHER IS BETTER; “GAP” IS THE LINEARITY TAX AGAINST EACH MODEL’S OWN ENCODE–DECODE REFERENCE (LOWER IS BETTER). “+ $f(0)$ ” ROWS ADD THE ENCODED SILENCE TO THE SUBTRACTED LATENT (ORIGIN CORRECTION). AGGREGATES OVER 5,880 UNITS FROM 49 TRACKS.

Not domain-driven. The $\text{SDR}_{\text{lin}}^{\text{gt}}$ improvement appears on all three test domains at both rates (+1.2 to +3.2 dB).

C. Downstream: latent-arithmetic stem removal

We remove a stem by latent subtraction and measure SI-SDR against the true residual on the MUSDB18 test set (49 of 50 tracks; one is excluded because its stems sum to the mixture at only 18.5 dB). This is an *oracle* setting: the removed stem is available and encoded, so it tests arithmetic, not blind separation. Subtraction uses coefficients (1, −1), which Eq. 1 does not cover, so we also report the origin-corrected form $g(f(x_{\text{mix}}) - f(x_{\text{stem}}) + f(0))$, which for an affine encoder decodes exactly $g(f(x_{\text{res}}))$, the model’s own reconstruction of the residual; the affine assumption is what makes the recovered *latent* exact. The shortfall against that encode–decode reference is the *linearity tax*.

Raw subtraction favours the mixing loss; origin correction substantially reduces the difference. Raw, $\mathcal{L}_{\text{mix}}^{\text{dec}}$ raises subtraction quality at 7.66× (+3.4→+5.8 dB) and the matched 15.3× pair goes from +1.2 to +5.2 dB. Adding $f(0)$ gains +2.0 to +6.0 dB for every evaluated GAN configuration, leaving each within about 0.5 dB of that reference in aggregate (Table III); per track the tail is wider. For these autoencoders, origin correction removes most of the mean subtraction advantage of mixing supervision: $\mathcal{L}_{\text{mix}}^{\text{dec}}$ versus no mixing is −0.03 dB (95% track-cluster CI [−0.15, +0.10]), the 15.3× pair +0.01 (−0.19, +0.21). Because these checkpoints were selected on this same split (Section III-C), we repeat the intervention on MoisesDB [14], 240 recordings absent from training and from selection, under a protocol registered before the audio was downloaded. Raw subtraction again favours the mixing models by +2.3 to +4.5 dB and recentering removes 87–95% of that, shrinking the gap on 229–236 of the 240 recordings, but the residue is resolvable there (+0.29, +0.12, +0.42 dB, intervals excluding zero). Convex mixing replicates on the same corpus: $\mathcal{L}_{\text{mix}}^{\text{dec}}$ gains +1.73 dB [+1.69, +1.78]

at 7.66× and the combined recipe +1.33 [+1.27, +1.39] at 15.3× for a reconstruction change of +0.04 and −0.07 dB, while the adversarial term *costs* 0.49 dB [−0.51, −0.48] of ground-truth equivariance and adds +4.06 dB of decode-vs-decode. For the held-out subtraction only, units are scored where the removed stem is active (24,648 of 28,800); the rule excludes silent *targets* and does not make the SI-SDR reference nonsilent, that reference being the remaining mixture. The three residual contrasts are $\mathcal{L}_{\text{mix}}^{\text{dec}}$ and $\mathcal{L}_{\text{mix}}^{\text{dec}}$ +disc against no mixing at 7.66×, and the combined recipe against no mixing at 15.3×. The comparator figures below return to the MUSDB18 test set. The no-mixing models also carry the larger offset ($\|f(0)\|$ 6.6 versus 5.4 at 7.66×), though offset norm alone does not determine the decoded error. What the loss buys is convex equivariance (Table I), which the correction leaves intact. The residual 0.4–0.5 dB is the arithmetic gap left after decoding, not a direct measure of non-affinity.

Baselines without latent arithmetic. The unchanged mixture scores +17.1 dB against the residual, above every model’s reconstruction reference, since the removed stem is a small share of the energy: the metric certifies arithmetic only relative to the model’s own reconstruction. Origin-corrected latent subtraction beats decode-then-subtract by +0.7 to +0.9 dB (CIs exclude zero); raw subtraction loses to it everywhere.

Intervals come from a paired cluster bootstrap over recordings (10^4 resamples), 49 tracks on the test set and 240 on the held-out corpus; mixing pairs join two chunks, about 80% of them from the same recording, and are weighted by the multiplicity of each member. They cover evaluation variance at the tested initialization only.

D. Equivariance across mixing ratios

$\mathcal{L}_{\text{mix}}^{\text{dec}}$ and its matched control are the pair swept under the same protocol (source-balanced, 1,200 samples per coefficient). The $\text{SDR}_{\text{lin}}^{\text{gt}}$ swing across α is 1.9 dB with $\mathcal{L}_{\text{mix}}^{\text{dec}}$ against 3.2 dB without, and $\mathcal{L}_{\text{mix}}^{\text{dec}}$ leads at all five coefficients (by 0.6 to 1.9 dB), so the property is not specific to $\alpha=0.5$. The $\mathcal{L}_{\text{mix}}^{\text{dec}}$ +disc sweep used the full test set, a different population, so we do not compare it here.

V. CONCLUSION

A controlled study identifies the decode-mixing loss, not an adversarial term on the mixed path, as the signal behind convex mixing-equivariance, with reconstruction SI-SDR similar across compression rates. Raw stem subtraction is instead governed largely by the latent offset: adding encoded silence brings subtraction in aggregate within about 1 dB of each model’s reconstruction reference and reduces the loss’s apparent advantage by 87–95% on a held-out corpus, exactly only in the affine limit. On a consistency-model autoencoder both latent geometry and ground-truth mixing improve while staying far from usable waveform arithmetic. **Limitations.** Comparisons are matched at one initialization, so each factor’s effect is identified there and not across runs; checkpoint selection used the test split (Section III-C); audio is mono and untested by listeners.

REFERENCES

- [1] H. Liu, Z. Chen, Y. Yuan, X. Mei, X. Liu, D. Mandic, W. Wang, and M. D. Plumbley, “AudioLDM: Text-to-audio generation with latent diffusion models,” in *Proc. of the 40th Int. Conf. on Machine Learning (ICML)*, 2023.
- [2] Z. Evans, J. D. Parker, C. Carr, Z. Zukowski, J. Taylor, and J. Pons, “Stable audio open,” in *Proc. of the IEEE Int. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, 2025.
- [3] R. Kumar, P. Seetharaman, A. Luebs, I. Kumar, and K. Kumar, “High-fidelity audio compression with improved RVQGAN,” in *Proc. of the 37th Int. Conf. on Neural Information Processing Systems (NeurIPS)*, 2023.
- [4] A. Défossez, J. Copet, G. Synnaeve, and Y. Adi, “High fidelity neural audio compression,” *Trans. on Machine Learning Research*, 2023.
- [5] B. Torres, M. Moussallam, and G. Meseguer-Brocal, “Learning linearity in audio consistency autoencoders via implicit regularization,” in *Proc. of the IEEE Int. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, 2026.
- [6] H. Zhang, M. Cisse, Y. N. Dauphin, and D. Lopez-Paz, “mixup: Beyond empirical risk minimization,” in *Proc. of the 6th Int. Conf. on Learning Representations (ICLR)*, 2018.
- [7] N. Zeghidour, A. Luebs, A. Omran, J. Skoglund, and M. Tagliasacchi, “SoundStream: An end-to-end neural audio codec,” *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, vol. 30, pp. 2709–2723, 2022.
- [8] J. Kong, J. Kim, and J. Bae, “HiFi-GAN: Generative adversarial networks for efficient and high fidelity speech synthesis,” in *Proc. of the 34th Int. Conf. on Neural Information Processing Systems (NeurIPS)*, 2020.
- [9] A. Caillon and P. Esling, “RAVE: A variational autoencoder for fast and high-quality neural audio synthesis,” *arXiv preprint arXiv:2111.05011*, 2021.
- [10] I. Higgins, L. Matthey, A. Pal, C. Burgess, X. Glorot, M. Botvinick, S. Mohamed, and A. Lerchner, “ β -VAE: Learning basic visual concepts with a constrained variational framework,” in *Proc. of the 5th Int. Conf. on Learning Representations (ICLR)*, 2017.
- [11] I. Higgins, D. Amos, D. Pfau, S. Racanière, L. Matthey, D. Rezende, and A. Lerchner, “Towards a definition of disentangled representations,” *arXiv preprint arXiv:1812.02230*, 2018.
- [12] M. Pasini, S. Lattner, and G. Fazekas, “Music2Latent: Consistency autoencoders for latent audio compression,” in *Proc. of the 25th Int. Society for Music Information Retrieval Conf. (ISMIR)*, 2024, pp. 111–119.
- [13] Z. Rafi, A. Liutkus, F.-R. Stöter, S. I. Mimilakis, and R. Bittner, “MUSDB18-HQ: An uncompressed version of MUSDB18,” <https://zenodo.org/record/3338373>, 2019.
- [14] I. Pereira, F. Araújo, F. Korzeniowski, and R. Vogl, “MoisesDB: A dataset for source separation beyond 4-stems,” in *Proc. of the 24th Int. Society for Music Information Retrieval Conf. (ISMIR)*, 2023, pp. 619–626.